

The Name of the Title Is Hope

Anonymous Author(s)

Abstract

LLM agents with web access are increasingly becoming the default for user-facing question answering. Reliable evaluation in this setting requires tests that remain unseen at train time and cannot be answered via verbatim lookup at inference. However, standard practice still relies on static benchmarks that age quickly and are vulnerable to contamination. We analyze two channels: pretraining leakage, where test items appear in the pretraining data, and retrieval-time leakage, where the system surfaces the exact target answer verbatim during evaluation and reports it rather than demonstrating genuine problem solving. We uncover evidence across five QA datasets for both channels and show that static evaluation can overestimate capability. We introduce DynamicWebQA, a dynamic evaluation framework for web-RAG that constructs questions and verifiable answers at evaluation time from current web sources. The framework enforces multi-document grounding, records evidence chains, and offers controllable complexity via source-set size, number of hops, and reasoning constraints. By generating fresh test instances, DynamicWebQA reduces memorization risk, mitigates retrieval-time leakage, and naturally covers time-sensitive queries.

CCS Concepts

• **Do Not Use This Code** → **Generate the Correct Terms for Your Paper**; *Generate the Correct Terms for Your Paper*; Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper.

Keywords

Do, Not, Use, This, Code, Put, the, Correct, Terms, for, Your, Paper

ACM Reference Format:

Anonymous Author(s). 2018. The Name of the Title Is Hope. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 10 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

Benchmark leakage refers to the contamination that occurs when a model has prior exposure to a benchmark’s test items during pretraining, leading to inflated evaluation scores and undermining the benchmark’s validity [46, 47]. Since most LLMs are pretrained indiscriminately on publicly available web content, there is a high likelihood that they might have encountered benchmark datasets

hosted on platforms like HuggingFace, GitHub, Kaggle, and similar sites. If the base LLM used in a WebQA agent has already seen the test examples during pretraining, the evaluation effectively becomes a case of testing on the training set — severely compromising the credibility of the reported results.

Dynamic benchmarking is an emerging evaluation paradigm that generates or adapts test samples at evaluation time, so models are scored on fresh, unseen data and testset contamination is avoided [19, 26, 34, 45, 47, 48]. That means the agent does not see the same test sample twice, effectively addressing issues such as LLM memorization.

In this work, we refer to information that can change over time, such as the stock price of Tesla, as evolving information. Arguably, this type of content could be labeled as dynamic content or temporal content, but that might cause confusion, since we are already using the term “dynamic” to describe an evaluation paradigm (dynamic benchmarking), and “temporal” might lead readers to associate it with time-series data. To avoid such confusion, we adopt the term evolving information to specifically denote factual content that may shift over time without implying a particular format or evaluation framework.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference acronym 'XX, Woodstock, NY

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2018/06

<https://doi.org/XXXXXXX.XXXXXXX>

2 Related Work

Open-Domain QA (ODQA). The ODQA task studies how to find and use evidence from large corpora to answer free-form questions. A sequence of datasets marked key shifts in task design: SQuAD [32] established span-based reading comprehension; Natural Questions [20] moved to real search queries; TriviaQA [16] stressed compositional, trivia-style questions with substantial lexical and syntactic variation; and HotpotQA [43] formalized multi-hop reasoning across multiple documents with sentence-level supporting facts and comparison questions, while including distractors to test focus under noise. Models such as fusion-in-decoder (FiD) [14] showed that aggregating many retrieved passages is a powerful way to synthesize scattered evidence.

Another important dimension in ODQA is *time*. RealTime QA [18] and FreshQA [39] evaluate whether systems can answer what is true *right now*, revealing failure modes when indices are stale or retrieved content is outdated or contradictory. These settings highlight needs that go beyond span extraction: recognizing when the answer has changed, when a question carries a false premise, and when recent evidence is required over older but higher-ranked documents. Notably, surface simplicity does not imply task simplicity. SimpleQA [41] demonstrates that short, seemingly straightforward questions can still require precise world knowledge or multi-step reasoning, indicating that difficulty is not well represented by length or phrasing.

Early ODQA followed a retrieve-then-read paradigm with extractive readers over IR-retrieved text. DrQA used TF-IDF/BM25 over Wikipedia coupled with a neural span extractor [7]. BERTserini then paired BM25 with a BERT reader to push end-to-end accuracy [42]. To support multi-hop questions, subsequent work leveraged structured signals. Asai et al. [3] learned to retrieve reasoning paths over the Wikipedia hyperlink graph to connect supporting pages. Sun et al. [35] introduced document–entity heterogeneous graphs (e.g., GraftNet), which aggregated evidence across passages and entities for open-domain QA. Dense retrieval reduced lexical mismatch: ORQA learned a latent retriever from QA pairs [22], and DPR’s dual-encoder with supervised negatives yielded large gains over strong BM25 baselines [17]. These systems still relied on static indices, motivating later generation-conditioned approaches like RAG.

Retrieval-Augmented Generation (RAG). Modern RAG approaches follow a common pipeline – given a query, retrieve top-*k* passages from a non-parametric memory and condition a generator on them. Foundational formulations include RAG [23], REALM [12], RETRO [4], and ATLAS [15], which together demonstrate that query-time retrieval improves knowledge-intensive tasks while enabling the external memory to be updated without re-training the parametric model. Recent surveys study components and variants of this paradigm [11].

In the literature, RAG is distinguished from broader tool-augmented generation – RAG requires *query-conditioned retrieval* whose contents directly condition the generator. Calculators, code execution, or static always-prepended context are typically treated as non-retrieval tools [27]. In practice, the retrieval index can be the open web or enterprise/local corpora, and the mechanism can be lexical

(e.g., BM25) or dense; the generator then aggregates evidence across multiple passages [14, 23].

Beyond single-shot RAG, agentic frameworks let models iteratively issue searches, follow links, and cite sources. WebGPT augments an LLM with browsing and requires collecting references during answering [29]. ReAct further couples reasoning traces with tool calls (e.g., search, calculator) in a unified prompting loop, enabling multi-hop planning and verification [44]. These approaches often embed RAG steps inside a broader plan-act-observe loop, allowing complex queries that require cross-document synthesis.

Data Contamination. Data contamination arises when evaluation items (or close paraphrases and answer-bearing spans) leak into the corpora used for pretraining, instruction tuning, or fine-tuning. In such cases, test behaviour can devolve into *lookup* of memorized strings rather than the intended inference or problem solving, inflating reported performance and obscuring brittleness [24, 33]. Prior studies connect leakage to memorization dynamics: larger models memorize more and earlier in training; duplication and longer prompts amplify verbatim regurgitation; and model capacity allows substantial unintended storage of datapoint-specific information [6, 28, 38]. At the same time, verbatim completion is not definitive evidence of membership – models can perfectly complete sequences never explicitly seen during training – implying membership tests must be interpreted cautiously [25]. The measurable impact of contamination is further mediated by dataset construction, model scale, and properties of pretraining data [24].

Ravaut et al. [33] organize detection techniques by *data availability*. With candidate corpora available (*open-data*), methods compare evaluation items against sources via instance-/dataset-level span matching, embedding similarity, or paraphrase detection. When training data are unavailable (*closed-data*), researchers rely on indirect evidence such as performance analyses, membership-inference attacks, and instance-level probes of memorization and model confidence. Despite blind spots, approaches based on span matching remain the practical default in most foundational efforts [5, 8, 10, 21, 30, 31] – as they are auditable and more importantly, tractable at web scale. Auditing a single question or sentence against a large crawl is a needle-in-a-haystack search; doing so for entire benchmarks across billions of documents renders more sophisticated but compute-heavy detectors infeasible. Moreover, behavioural signals alone rarely certify contamination [24, 25].

Retrieval-augmented and agentic QA introduce a new concern – these systems may retrieve and quote target items from web or enterprise indices at *run time*, creating search-time data contamination that boosts scores without demonstrating reasoning [13]. Han et al. [13] show that even small leakage rates (e.g., on the order of a few percent) can reorder competitive leaderboards once online access is enabled, undermining trust relative to purely offline evaluation.

Dynamic Benchmarking. Most benchmarks for agentic or retrieval-augmented QA remain static and publicly accessible (e.g., HotpotQA), which makes them susceptible to contamination and memorization [6, 24, 28, 33, 43]. As models and training corpora grow, leakage can inflate reported scores on static benchmarks. This has motivated *dynamic* evaluation paradigms that are more robust towards contamination and memorization [9, 34, 45, 49].

Early dynamic platforms such as Dynabench and Dynaboard introduced human- and model-in-the-loop data collection [19, 26, 37], but their reliance on crowdsourcing limits scalability. More recent approaches leverage domain-specific settings (e.g., AndroidWorld for interactive agents [34]) or graph-based generation to create evaluation samples with controllable complexity and cross-round consistency [9, 45, 49]. While graph-anchored methods have proven effective in tasks like reasoning and KGQA [9, 45], their application to ODQA is still limited, a gap we address in this work.

3 Modeling Contamination in RAG-Based QA

RAG systems combine a parametric model M with external sources retrieved at inference. This creates two leakage channels: (i) *pre-training contamination*, where an evaluation item (or a close paraphrase / answer-bearing span) appears in the corpus used to pre-train M ; and (ii) *retrieval-time contamination*, where the retriever surfaces such material during evaluation, enabling lookup rather than reasoning. Our experiments instantiate web-enabled RAG “Web-RAG”, but the formalism below applies to RAG broadly.

3.1 Definitions

Let M be pretrained on corpus D_M . Let D_{eval} contain items $x = (q, a)$ with question q and answer a . Let $\mathcal{R}(q)$ be the set of documents returned by the retriever for query q . Following Ravaut et al. [33], let f be a similarity function and τ a threshold; $\tau=1$ corresponds to *strict/verbatim* contamination.

We study both *question-only* and *question+answer* leakage, noting that even input-only (question-only) leakage can shift measured performance [24]. We consider two views of an item:

$$\pi_{\text{in}}(q, a) = q \quad (\text{question-only, a.k.a. input-only}),$$

$$\pi_{\text{in+lab}}(q, a) = q \parallel a \quad (\text{question+answer, a.k.a. input+label}).$$

Contamination criterion. For a channel $C \in \{\text{pre}, \text{ret}\}$ and view $v \in \{\text{in}, \text{in+lab}\}$, we say D_{eval} is contaminated if

$$\exists x \in D_{\text{eval}}, \exists d \in S_C(x) \text{ s.t. } f(\pi_v(x), d) \geq \tau, \quad (1)$$

with source sets

$$S_{\text{pre}}(x) = D_M, \quad S_{\text{ret}}(x) = \mathcal{R}(q).$$

When a contaminant must be explicitly answer-bearing, we additionally require a predicate $g(a, d) \in \{0, 1\}$ (e.g., answer mention or high-overlap), and record positives of

$$\exists x, d \text{ s.t. } f(q, d) \geq \tau_q \wedge g(a, d) = 1. \quad (2)$$

3.2 Implementation details

We instantiate f using five auditable open-data detectors frequently used in LLM pre-training studies and consolidated in [33]:

- **Exact Match (EM)** [10].
- **8-gram words** (GPT-2 pre-training study) [31].
- **13-gram words** (GPT-3 pre-training study) [5].
- **50-character substring** (GPT-4 technical report) [30].
- **8-gram words with 70% threshold** (PaLM pre-training study) [8].

Span-based detectors scale to long documents, handle the length mismatch between short questions and web pages, and yield concrete evidence spans. While semantic detectors (embeddings / paraphrase) may increase recall, they are harder to audit and become infeasible across disparate lengths; we therefore adopt span-based methods as primary signals, consistent with cautions in [25] and common audit practice [33].

Pretraining channel (candidate retrieval). Full membership tests against all datapoints in web-scale pretraining corpora are infeasible. Following candidate-set approximations in open-data audits [24], we index the Common Crawl C4 corpus in an Elasticsearch-compatible service (AWS OpenSearch). For each $x = (q, a)$, we retrieve top- k candidates with BM25 (default $k=1000$) and compute contamination using Eq. (1) for *both views* (π_{in} and $\pi_{\text{in+lab}}$), as well as the *answer-bearing* variant Eq. (2) (question view with $g(a, d)$). Any positive flag indicates potential pretraining contamination. This candidate pool serves as $S_{\text{pre}}(x)$.

Retrieval-time channel (web-RAG). For each q , we query SearxNG [1] with the raw question string, pool results from the first three results pages across underlying engines, and define $\mathcal{R}(q)$ accordingly. We compute contamination over this pool using Eq. (1) for *both views* and Eq. (2) (answer-bearing, question view), and we also use the same pool as retrieval context for answer generation. This isolates run-time leakage introduced by search from memorization in M . Here $S_{\text{ret}}(x) = \mathcal{R}(q)$.

4 DYNAMICWEBQA Framework

We propose a method that leverages *seed graphs* to anchor future question generation while avoiding fixed test items that are easy to memorize or contaminate. At each evaluation round, every graph yields a fresh QA pair produced from the same underlying anchors, so rounds are semantically comparable without reusing static instances.

4.1 Seed Graphs

At evaluation round t , we maintain a set $\mathcal{S}_t = \{G_1^{(t)}, \dots, G_{n_t}^{(t)}\}$ of seed graphs. Each graph $G_i^{(t)} = (V_i^{(t)}, E_i^{(t)}, \Xi_i^{(t)})$ contains: (i) a node set $V_i^{(t)}$ of *documents* retrieved at evaluation time, (ii) an edge set $E_i^{(t)}$ that records how documents have been jointly used to form past questions, and (iii) metadata $\Xi_i^{(t)}$ (timestamps, operator labels, usage counts). We initialize each graph with document nodes. As new QA items are generated from the graph, we *add edges* between the documents used in that item and label those edges with the applied operator and round. This builds a usage history over time, supports collision control, and keeps nodes aligned with current web sources when refreshed at later rounds.

4.2 Claim Extraction

Given the current document nodes $V_i^{(t)}$, we extract one-sentence, verifiable claims and their supporting spans using a fixed LLM prompt (“Claim Extraction Prompt”). Let

$$C_i^{(t)} = \bigcup_{d \in V_i^{(t)}} \{(d, \text{claim}, \text{span})\}$$

denote the set of extracted claims, bucketed by their source document d . This step yields atomic, referenceable units that we later compose into questions.

4.3 Operator-Controlled QA Composition

Question complexity and the type of reasoning required are jointly controlled by (i) a small set of composition operators $\mathcal{O}=\{\text{conjunction, temporal, comparison}\}$ and (ii) the hop count h – the number of distinct document buckets used – constrained to $h \in [h_{\min}, h_{\max}]$.

Each operator is realized by a fixed prompt template (see Appendix): conjunction enforces a logical AND so the answer changes if any referenced claim were false; temporal asks for an ordering or interval over time; and comparison requires a contrast (e.g., higher/lower, earlier/later, ratio). We specify $h \in [h_{\min}, h_{\max}]$ to bound multi-hop complexity, with operator-specific constraints enforced by the templates.

4.4 QA Generation

The generation function Gen takes the graph, a random seed, and fixed configuration to produce one QA pair:

$$(q_{i,t}, a_{i,t}, U_{i,t}) = \text{Gen}(G_i^{(t)}; r_{i,t}, \theta), \quad (3)$$

where θ fixes the prompt templates and decoding settings (temperature, top- p , stop tokens), and $r_{i,t} \sim \mathcal{R}$ supplies per-sample stochasticity. Internally, Gen: (1) samples an operator $o \in \mathcal{O}$ and a hop count h within the configured range; (2) selects h document buckets from $C_i^{(t)}$ and assembles the operator-specific template with the selected claims; and (3) invokes an LLM to return a JSON object containing the list of used claims, the question, and a single concise answer. We write $U_{i,t}$ for the generation trace (operator o , documents used, and claim identifiers). Holding θ fixed and varying only $r_{i,t}$ produces i.i.d. samples per graph while keeping composition policy constant.

4.5 Graph Update and History

After generation, we update the graph to record usage:

$$G_i^{(t+1)} = \text{Update}(G_i^{(t)}, U_{i,t}),$$

which adds edges between all documents referenced in $U_{i,t}$ and annotates each edge with $(o, \text{round } t, \text{QA id})$. This history supports two practical goals: (i) tracking how documents have been combined across rounds and operators, and (ii) reducing collisions by allowing downstream filters to avoid reusing the same document combinations. The node set $V_i^{(t)}$ can also be refreshed to reflect changing sources.

4.6 Sampling and Distributional Properties

In round t , with the configuration θ held fixed (prompts, operator/hop ranges, decoding settings), the only randomness comes from per-graph seeds $r_{i,t} \sim \mathcal{R}$, which drive operator/hop selection, claim selection, and decoding inside Gen. For each graph, this induces a distribution over QA pairs:

$$\mathcal{P}_{i,t} = \text{Dist}(\text{Gen}(G_i^{(t)}; r_{i,t}, \theta)). \quad (4)$$

With independent seeds, the round- t dataset $D_t = \{(q_{1,t}, a_{1,t}), \dots, (q_{n_t,t}, a_{n_t,t})\}$ is an independent draw from the product measure

$$\mathcal{P}_t = \bigotimes_{i=1}^{n_t} \mathcal{P}_{i,t}. \quad (5)$$

For a model M and any per-question metric score $\text{score}((q, a), M)$ (e.g., EM, F1, log-loss), the aggregate dataset score is

$$\text{Agg}(D_t, M) = \frac{1}{n_t} \sum_{i=1}^{n_t} \text{score}((q_{i,t}, a_{i,t}), M). \quad (6)$$

By linearity of expectation,

$$\mathbb{E}_{D_t \sim \mathcal{P}_t} [\text{Agg}(D_t, M)] = \frac{1}{n_t} \sum_{i=1}^{n_t} \mathbb{E}_{(q,a) \sim \mathcal{P}_{i,t}} [\text{score}((q, a), M)], \quad (7)$$

so the expected aggregate is the average of the per-graph expectations, giving each seed graph equal weight.

4.7 Cross-Round Reporting and Compatibility

Across rounds, the benchmark maintains a persistent set of seed graphs $\mathcal{S}$. As QA items are generated, edges and usage history accumulate. The design remains compatible with evolving sources: nodes represent current documents, so when a source changes, claims can be re-extracted the next time that graph is sampled. Within intervals where the node set and configuration θ are unchanged, successive generations for a given graph are i.i.d.; over longer horizons, the per-graph distribution $\mathcal{P}_{i,t}$ can drift with source evolution. To report results under these conditions, we use two complementary views:

i) *Snapshot Evaluation*. Choose a reference round t^* and evaluate all systems on its realized dataset D_{t^*} . This fixes both the source set and the random draw, yielding identical test conditions.

ii) *Macro-Averaged Score*. For longitudinal tracking over a window $W = \{t_1, \dots, t_T\}$ (e.g., the most recent T rounds), compute

$$\text{Score}_{\text{macro}}(M) = \frac{1}{T} \sum_{j=1}^T \text{Agg}(D_{t_j}, M). \quad (8)$$

Each D_{t_j} is an i.i.d. draw from its round-specific $\mathcal{P}_{t_j}$, so this macro-average summarizes average effectiveness under natural benchmark evolution without letting any single round dominate.

4.8 Verifying Distributional Consistency

Under a stable generation policy (with θ held fixed), the embedding geometry should be invariant across rounds. Sentence embeddings are used as a practical proxy for topical relatedness, where proximity often reflects overlapping content or cues. Consequently, if the topic mix is stable, the distribution of pairwise distances should remain unchanged across rounds. Topic composition can change in ways that leave the average embedding unchanged—for example, similar topics in different proportions or broader within-topic spread—yet these changes alter the pattern of pairwise distances. Round-wise differences are therefore assessed using a distance-based permutation MANOVA (PERMANOVA) [2] as the omnibus test, and pairwise differences using the energy two-sample test [36].

Table 1: Omnibus PERMANOVA across $K=10$ rounds (Euclidean on ℓ_2 -normalized embeddings; 10k permutations).

Split	N	Pseudo- F	R^2 (%)	Perms	p -value
Highly Contaminated	5,000	0.7429	0.1338	10,000	≥ 0.99
Randomly Sampled	4,996	0.6476	0.1168	10,000	≥ 0.99

Embedding and distances. Each QA pair is concatenated and embedded with a sentence embedding model (intfloat/e5-base-v2 [40] in our experiments). Embeddings are ℓ_2 -normalized, so Euclidean and cosine distances are monotone-equivalent. Let D denote the Euclidean distance matrix over all items.

Omnibus test across K rounds. PERMANOVA is applied to D with round labels (10,000 permutations). Reported statistics are the pseudo- F , permutation p -value, and R^2 . This serves as a K -sample omnibus test for differences in location and dispersion in the distance space.

Pairwise localization. For all $\binom{K}{2}$ round pairs, the energy two-sample test is applied on D (10,000 permutations) with Benjamini–Hochber FDR control at $\alpha=0.05$. Reported quantities are the E-statistic and permutation p -value.

Descriptive visuals. (i) A 2D PCA of embeddings colored by round (descriptive intermixing). (ii) An energy–distance heat map summarizing cross-round separation after accounting for within-round dispersion; larger off-diagonal values indicate greater separation.

5 Results

5.1 Contamination Modeling in QA Datasets




Figure 1: Contamination rates by dataset across the two leakage channels defined in the modeling framework: (a) pre-training contamination (source set $S_{pre}=D_M$), and (b) retrieval channel contamination (source set $S_{ret}=\mathcal{R}(q)$).

5.2 Model Performance

5.3 Distributional Consistency: Results

Omnibus. PERMANOVA detects no round effect in either split (Table 1). With 10,000 permutations the p -value resolution is 10^{-4} ; both splits yield $p \geq 0.99$ and $R^2 \leq 0.134\%$, indicating negligible variance explained by “round”.




Figure 2: Top contamination sources during retrieval for each dataset. For each dataset, we show the top-5 domains by number of contaminated instances.

Pairwise. For the 45 round pairs per split, no comparison is significant after BH–FDR ($\alpha=0.05$). E-statistics are tiny (10^{-4} – 10^{-3}) and centered near zero (Table 2); maxima are $< 10^{-3}$.

Visual evidence. PCA projections (Fig. 5) show full intermixing by round (PC1/PC2 $< 4\%$ explained each). Drift heatmaps of signed




Figure 3: Rank distribution of contaminated evidence surfaced by the retriever, shown per dataset. Histograms separately show question-only and question+answer contamination where available.






Figure 4: Caption TBD

Table 2: Summary of pairwise Energy tests (signed, U-centered E-statistic).


Split	#Sig (FDR)	Mean	Median
Highly Contaminated	0/45	-5.14×10^{-4}	-6.31×10^{-4}
Randomly Sampled	0/45	-6.58×10^{-4}	-7.10×10^{-4}

**Figure 5: PCA projections of ℓ_2 -normalized embeddings for Highly Contaminated (left) and Randomly Sampled (right) splits. Points are colored by round; both splits show full intermixing with PC1/PC2 each explaining $<4\%$ variance.****Figure 6: Pairwise drift heatmaps. Left: Highly Contaminated; Right: Randomly Sampled. Cells show the E -statistic from the Energy two-sample test (hyppo.ksample.Energy) computed on Euclidean distances of ℓ_2 -normalized embeddings; 10k permutations with BH-FDR over 45 round pairs.**

E-statistics (Fig. 6) concentrate near zero without structure. Centered mean-distance maps and relative deviations (Fig. 7) show max absolute effects $\leq 1.5 \times 10^{-3}$, i.e., $<0.25\%$ of the within-round scale.

6 Conclusion

TBD

**Figure 7: Centered mean distance across rounds. Left: Highly Contaminated; Right: Randomly Sampled. Each cell shows cross-round mean distance minus the average within-round distance (proportional to the biased Energy distance). Values cluster near zero with maxima $\leq 1.7 \times 10^{-3}$.**

References

- [1] 2025. Searxng/Searxng. SearXNG.
- [2] Marti J. Anderson. 2001. A New Method for Non-Parametric Multivariate Analysis of Variance. *Austral Ecology* 26, 1 (2001), 32–46. doi:10.1111/j.1442-9993.2001.01070.pp.x
- [3] Akari Asai, Kazuma Hashimoto, Hannaneh Hajishirzi, Richard Socher, and Caiming Xiong. 2020. Learning to Retrieve Reasoning Paths over Wikipedia Graph for Question Answering. arXiv:1911.10470 [cs] doi:10.48550/arXiv.1911.10470 Comment: Published as a conference paper at ICLR 2020. Code is available at https://github.com/AkariAsai/learning_to_retrieve_reasoning_paths.
- [4] Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, Diego de Las Casas, Aurelia Guy, Jacob Menick, Roman Ring, Tom Hennigan, Saffron Huang, Loren Maggiore, Chris Jones, Albin Cassirer, Andy Brock, Michela Paganini, Geoffrey Irving, Oriol Vinyals, Simon Osindero, Karen Simonyan, Jack W. Rae, Erich Elsen, and Laurent Sifre. 2022. Improving Language Models by Retrieving from Trillions of Tokens. arXiv:2112.04426 [cs] doi:10.48550/arXiv.2112.04426 Comment: Fix incorrect reported numbers in Table 14.
- [5] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Aidan Clark, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language Models Are Few-Shot Learners. arXiv:2005.14165 [cs] doi:10.48550/arXiv.2005.14165 Comment: 40+32 pages.
- [6] Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan Zhang. 2023. Quantifying Memorization Across Neural Language Models. arXiv:2202.07646 [cs] doi:10.48550/arXiv.2202.07646
- [7] Danqi Chen, Adam Fisch, Jason Weston, and Antoine Bordes. 2017. Reading Wikipedia to Answer Open-Domain Questions. arXiv:1704.00051 [cs] doi:10.48550/arXiv.1704.00051 Comment: ACL2017, 10 pages.
- [8] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, Parker Schuh, Kensen Shi, Sasha Tsvyashchenko, Joshua Maynez, Abhishek Rao, Parker Barnes, Yi Tay, Noam Shazeer, Vinodkumar Prabhakaran, Emily Reif, Nan Du, Ben Hutchinson, Reiner Pope, James Bradbury, Jacob Austin, Michael Isard, Guy Gur-Ari, Pengcheng Yin, Toju Duke, Anselm Levskaya, Sanjay Ghemawat, Sunipa Dev, Henryk Michalewski, Xavier Garcia, Vedant Misra, Kevin Robinson, Liam Fedus, Denny Zhou, Daphne Ippolito, David Luan, Hyeontaek Lim, Barret Zoph, Alexander Spiridonov, Ryan Sepassi, David Dohan, Shivani Agrawal, Mark Omernick, Andrew M. Dai, Thanumalayan Sankaranarayanan Pillai, Marie Pellat, Aitor Lewkowycz, Erica Moreira, Rewon Child, Oleksandr Polozov, Katherine Lee, Zongwei Zhou, Xuezhi Wang, Brennan Saeta, Mark Diaz, Orhan Firat, Michele Catasta, Jason Wei, Kathy Meier-Hellstern, Douglas Eck, Jeff Dean, Slav Petrov, and Noah Fiedel. 2022. PaLM: Scaling Language Modeling with Pathways. arXiv:2204.02311 [cs] doi:10.48550/arXiv.2204.02311
- [9] Preetam Prabhur Srikar Dammu, Himanshu Naidu, and Chirag Shah. 2025. Dynamic-KGQA: A Scalable Framework for Generating Adaptive Question Answering Datasets. In *Proceedings of the 48th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '25)*. Association for Computing Machinery, New York, NY, USA, 3498–3508. doi:10.1145/3726302.3730324
- [10] Jesse Dodge, Maarten Sap, Ana Marasović, William Agnew, Gabriel Ilharco, Dirk Groeneveld, Margaret Mitchell, and Matt Gardner. 2021. Documenting Large Webtext Corpora: A Case Study on the Colossal Clean Crawled Corpus. arXiv:2104.08758 [cs] doi:10.48550/arXiv.2104.08758 Comment: EMNLP 2021 accepted paper camera ready version.
- [11] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofeng Wang. 2024. Retrieval-Augmented Generation for Large Language Models: A Survey. arXiv:2312.10997 [cs] doi:10.48550/arXiv.2312.10997 Comment: Ongoing Work.
- [12] Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Mingwei Chang. 2020. Retrieval Augmented Language Model Pre-Training. In *Proceedings of the 37th International Conference on Machine Learning*. PMLR, 3929–3938.
- [13] Ziwen Han, Meher Mankikar, Julian Michael, and Zifan Wang. 2025. Search-Time Data Contamination. arXiv:2508.13180 [cs] doi:10.48550/arXiv.2508.13180
- [14] Gautier Izacard and Edouard Grave. 2021. Leveraging Passage Retrieval with Generative Models for Open Domain Question Answering. arXiv:2007.01282 [cs] doi:10.48550/arXiv.2007.01282
- [15] Gautier Izacard, Patrick Lewis, Maria Lomeli, Lucas Hosseini, Fabio Petroni, Timo Schick, Jane Dwivedi-Yu, Armand Joulin, Sebastian Riedel, and Edouard Grave. 2022. Atlas: Few-shot Learning with Retrieval Augmented Language Models. arXiv:2208.03299 [cs] doi:10.48550/arXiv.2208.03299
- [16] Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. 2017. TriviaQA: A Large Scale Distantly Supervised Challenge Dataset for Reading Comprehension. arXiv:1705.03551 [cs] doi:10.48550/arXiv.1705.03551 Comment: Added references, fixed typos, minor baseline update.
- [17] Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. Dense Passage Retrieval for Open-Domain Question Answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Bonnie Webber, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 6769–6781. doi:10.18653/v1/2020.emnlp-main.550
- [18] Jungo Kasai, Keisuke Sakaguchi, Yoichi Takahashi, Ronan Le Bras, Akari Asai, Xinyan Yu, Dragomir Radev, Noah A. Smith, Yejin Choi, and Kentaro Inui. 2024. RealTime QA: What's the Answer Right Now? arXiv:2207.13332 [cs] doi:10.48550/arXiv.2207.13332 Comment: RealTime QA Website: <https://realtimeqa.github.io/>.
- [19] Douwe Kiela, Max Bartolo, Yixin Nie, Divyansh Kaushik, Atticus Geiger, Zhengxuan Wu, Bertie Vidgen, Grusha Prasad, Amanpreet Singh, Pratik Ringshia, Zhiyi Ma, Tristan Thrush, Sebastian Riedel, Zeerak Waseem, Pontus Stenetorp, Robin Jia, Mohit Bansal, Christopher Potts, and Adina Williams. 2021. Dynabench: Rethinking Benchmarking in NLP. arXiv:2104.14337 [cs] doi:10.48550/arXiv.2104.14337 Comment: NAACL 2021.
- [20] Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. 2019. Natural Questions: A Benchmark for Question Answering Research. *Transactions of the Association for Computational Linguistics* 7 (2019), 452–466. doi:10.1162/tacl_a_00276
- [21] Ariel N. Lee, Cole J. Hunter, and Nataniel Ruiz. 2024. Platypus: Quick, Cheap, and Powerful Refinement of LLMs. arXiv:2308.07317 [cs] doi:10.48550/arXiv.2308.07317 Comment: Workshop on Instruction Tuning and Instruction Following at NeurIPS 2023.
- [22] Kenton Lee, Ming-Wei Chang, and Kristina Toutanova. 2019. Latent Retrieval for Weakly Supervised Open Domain Question Answering. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, Anna Korhonen, David Traum, and Lluís Màrquez (Eds.). Association for Computational Linguistics, Florence, Italy, 6086–6096. doi:10.18653/v1/P19-1612
- [23] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems*, Vol. 33. Curran Associates, Inc., 9459–9474.
- [24] Yucheng Li, Frank Guerin, and Chenghua Lin. 2024. An Open Source Data Contamination Report for Large Language Models. arXiv:2310.17589 [cs] doi:10.48550/arXiv.2310.17589
- [25] Ken Ziyu Liu, Christopher A. Choquette-Choo, Matthew Jagielski, Peter Kairouz, Sammi Koyejo, Percy Liang, and Nicolas Papernot. 2025. Language Models May Verbatim Complete Text They Were Not Explicitly Trained On. arXiv:2503.17514 [cs] doi:10.48550/arXiv.2503.17514 Comment: Main text: 9 pages, 7 figures, 1 table. Appendix: 29 pages, 20 tables, 15 figures.
- [26] Zhiyi Ma, Kawin Ethayarajh, Tristan Thrush, Somya Jain, Ledell Wu, Robin Jia, Christopher Potts, Adina Williams, and Douwe Kiela. 2021. Dynaboard: An Evaluation-As-A-Service Platform for Holistic Next-Generation Benchmarking. arXiv:2106.06052 [cs] doi:10.48550/arXiv.2106.06052
- [27] Grégoire Mialon, Roberto Dessi, Maria Lomeli, Christoforos Nalmpantis, Ram Pasunuru, Roberta Raileanu, Baptiste Rozière, Timo Schick, Jane Dwivedi-Yu, Asli Celikyilmaz, Edouard Grave, Yann LeCun, and Thomas Scialom. 2023. Augmented Language Models: A Survey. arXiv:2302.07842 [cs] doi:10.48550/arXiv.2302.07842
- [28] John X. Morris, Chawin Sitawarin, Chuan Guo, Narine Kokhlikyan, G. Edward Suh, Alexander M. Rush, Kamalika Chaudhuri, and Saeed Mahloujifar. 2025. How Much Do Language Models Memorize? arXiv:2505.24832 [cs] doi:10.48550/arXiv.2505.24832
- [29] Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, Xu Jiang, Karl Cobbe, Tyna Eloundou, Gretchen Krueger, Kevin Button, Matthew Knight, Benjamin Chess, and John Schulman. 2022. WebGPT: Browser-assisted Question-Answering with Human Feedback. arXiv:2112.09332 [cs] doi:10.48550/arXiv.2112.09332 Comment: 32 pages.
- [30] OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altmenschmidt, Sam Altman, Shyamal Anadkat, Red Avila, Igor Babuschkin, Suchir Balaji, Valerie Balcom, Paul Baltescu, Haiming Bao, Mohammad Bavarian, Jeff Belgum, Irwan Bello, Jake Berdine, Gabriel Bernadett-Shapiro, Christopher Berner, Lenny Boudeno, Oleg Boiko, Madelaine Boyd, Anna-Luisa Brakman, Greg Brockman, Tim Brooks, Miles Brundage, Kevin Button, Trevor Cai, Rosie Campbell, Andrew Cann, Brittany Carey, Chelsea Carlson, Rory Carmichael, Brooke Chan, Che Chang, Fotis Chantzis, Derek Chen, Sully Chen, Ruby Chen, Jason Chen, Mark Chen, Ben Chess, Chester Cho, Casey Chu, Hyung Won Chung, Dave Cummings, Jeremiah Currier, Yunxing Dai, Cory Decareaux, Thomas Degry, Noah Deutsch, Damien Deville, Arka Dhar, David Dohan, Steve Dowling, Sheila

- Dunning, Adrien Ecoffet, Atty Eleti, Tyna Eloundou, David Farhi, Liam Fedus, Niko Felix, Simón Posada Fishman, Juston Forte, Isabella Fulford, Leo Gao, Elie Georges, Christian Gibson, Vik Goel, Tarun Gogineni, Gabriel Goh, Rapha Gontijo-Lopes, Jonathan Gordon, Morgan Grafstein, Scott Gray, Ryan Greene, Joshua Gross, Shixiang Shane Gu, Yufei Guo, Chris Hallacy, Jesse Han, Jeff Harris, Yuchen He, Mike Heaton, Johannes Heidecke, Chris Hesse, Alan Hickey, Wade Hickey, Peter Hoeschele, Brandon Houghton, Kenny Hsu, Shengli Hu, Xin Hu, Joost Huizinga, Shantanu Jain, Shawn Jain, Joanne Jang, Angela Jiang, Roger Jiang, Haozhun Jin, Denny Jin, Shino Jomoto, Billie Jonn, Heewoo Jun, Tomer Kaftan, Lukasz Kaiser, Ali Kamali, Ingmar Kanitscheider, Nitish Shirish Keskar, Tabarak Khan, Logan Kilpatrick, Jong Wook Kim, Christina Kim, Yongjik Kim, Jan Hendrik Kirchner, Jamie Kiros, Matt Knight, Daniel Kokotajlo, Łukasz Kondraciuk, Andrew Kondrich, Aris Konstantinidis, Kyle Kosic, Gretchen Krueger, Vishal Kuo, Michael Lampe, Ikai Lan, Teddy Lee, Jan Leike, Jade Leung, Daniel Levy, Chak Ming Li, Rachel Lin, Molly Lin, Stephanie Lin, Mateusz Litwin, Theresa Lopez, Ryan Lowe, Patricia Lue, Anna Makanju, Kim Malfacini, Sam Manning, Todor Markov, Yaniv Markovski, Bianca Martin, Katie Mayer, Andrew Mayne, Bob McGrew, Scott Mayer McKinney, Christine McLeavey, Paul McMillan, Jake McNeil, David Medina, Aalok Mehta, Jacob Menick, Luke Metz, Andrey Mishchenko, Pamela Mishkin, Vinnie Monaco, Evan Morikawa, Daniel Mossing, Tong Mu, Mira Murati, Oleg Murk, David Mély, Ashvin Nair, Reiichiro Nakano, Rajeev Nayak, Arvind Neelakantan, Richard Ngo, Hyeonwoo Noh, Long Ouyang, Cullen O’Keefe, Jakub Pachocki, Alex Paino, Joe Palermo, Ashley Pantuliano, Giam Battista Parascandolo, Joel Parish, Emy Parparita, Alex Passos, Mikhail Pavlov, Andrew Peng, Adam Perelman, Filipe de Avila Belbute Peres, Michael Petrov, Henrique Ponde de Oliveira Pinto, Michael, Pokorny, Michelle Pokrass, Vitchyr H. Pong, Tolly Powell, Alethea Power, Boris Power, Elizabeth Proehl, Raul Puri, Alec Radford, Jack Rae, Aditya Ramesh, Cameron Raymond, Francis Real, Kendra Rimbach, Carl Ross, Bob Rotsted, Henri Roussez, Nick Ryder, Mario Saltarelli, Ted Sanders, Shibani Santurkar, Girish Sastry, Heather Schmidt, David Schnurr, John Schulman, Daniel Selsam, Kyla Sheppard, Toki Sherbakov, Jessica Shieh, Sarah Shoker, Pranav Shyam, Szymon Sidor, Eric Sigler, Maddie Simens, Jordan Sitkin, Katarina Slama, Ian Sohl, Benjamin Sokolowsky, Yang Song, Natalie Staudacher, Felipe Petroski Such, Natalie Summers, Ilya Sutskever, Jie Tang, Nikolas Tezak, Madeleine B. Thompson, Phil Tillet, Amin Tootoonchian, Elizabeth Tseng, Preston Tuggle, Nick Turley, Jerry Tworek, Juan Felipe Cerón Uribe, Andrea Vallone, Arun Vijayvergiya, Chelsea Voss, Carroll Wainwright, Justin Jay Wang, Alvin Wang, Ben Wang, Jonathan Ward, Jason Wei, C. J. Weinmann, Akila Welihinda, Peter Welinder, Jiayi Weng, Lilian Weng, Matt Wiethoff, Dave Willner, Clemens Winter, Samuel Wolrich, Hannah Wong, Lauren Workman, Sherwin Wu, Jeff Wu, Michael Wu, Kai Xiao, Tao Xu, Sarah Yoo, Kevin Yu, Qiming Yuan, Wojciech Zaremba, Rowan Zellers, Chong Zhang, Marvin Zhang, Shengjia Zhao, Tianhao Zheng, Juntang Zhuang, William Zhuk, and Barret Zoph. 2024. GPT-4 Technical Report. arXiv:2303.08774 [cs] doi:10.48550/arXiv.2303.08774 Comment: 100 pages; updated authors list; fixed author names and added citation.
- [31] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. [n.d.]. Language Models Are Unsupervised Multitask Learners. ([n.d.]).
- [32] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. SQuAD: 100,000+ Questions for Machine Comprehension of Text. arXiv:1606.05250 [cs] doi:10.48550/arXiv.1606.05250 Comment: To appear in Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing (EMNLP).
- [33] Mathieu Ravaut, Bosheng Ding, Fangkai Jiao, Hailin Chen, Xingxuan Li, Ruochen Zhao, Chengwei Qin, Caiming Xiong, and Shafiq Joty. 2025. A Comprehensive Survey of Contamination Detection Methods in Large Language Models. arXiv:2404.00699 [cs] doi:10.48550/arXiv.2404.00699 Comment: Accepted by TMLR in July 2025. 18 pages, 1 figure, 3 tables.
- [34] Christopher Rawles, Sarah Clinckemaitie, Yifan Chang, Jonathan Waltz, Gabrielle Lau, Marybeth Fair, Alice Li, William Bishop, Wei Li, Folawiyi Campbell-Ajala, Daniel Toyama, Robert Berry, Divya Tyamagundlu, Timothy Lillcrap, and Oriana Riva. 2025. AndroidWorld: A Dynamic Benchmarking Environment for Autonomous Agents. arXiv:2405.14573 [cs] doi:10.48550/arXiv.2405.14573
- [35] Haitian Sun, Bhuwan Dhingra, Manzil Zaheer, Kathryn Mazaitis, Ruslan Salakhutdinov, and William W. Cohen. 2018. Open Domain Question Answering Using Early Fusion of Knowledge Bases and Text. arXiv:1809.00782 [cs] doi:10.48550/arXiv.1809.00782 Comment: EMNLP 2018.
- [36] Gábor J. Székely and Maria L. Rizzo. 2013. Energy Statistics: A Class of Statistics Based on Distances. *Journal of Statistical Planning and Inference* 143, 8 (Aug. 2013), 1249–1272. doi:10.1016/j.jspi.2013.03.018
- [37] Tristan Thrush, Kushal Tirumala, Anmol Gupta, Max Bartolo, Pedro Rodriguez, Tariq Kane, William Gaviria Rojas, Peter Mattson, Adina Williams, and Douwe Kiela. 2022. Dynatask: A Framework for Creating Dynamic AI Benchmark Tasks. arXiv:2204.01906 [cs] doi:10.48550/arXiv.2204.01906 Comment: ACL System Demos 2022.
- [38] Kushal Tirumala, Aram H. Markosyan, Luke Zettlemoyer, and Armen Aghajanyan. 2022. Memorization Without Overfitting: Analyzing the Training Dynamics of Large Language Models. arXiv:2205.10770 [cs] doi:10.48550/arXiv.2205.10770
- [39] Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan Sung, Denny Zhou, Quoc Le, and Thang Luong. 2023. FreshLLMs: Refreshing Large Language Models with Search Engine Augmentation. arXiv:2310.03214 [cs] doi:10.48550/arXiv.2310.03214 Comment: Preprint, 26 pages, 10 figures, 5 tables; Added FreshEval.
- [40] Liang Wang, Nan Yang, Xiaolong Huang, Binxing Jiao, Linjun Yang, Daxin Jiang, Rangan Majumder, and Furu Wei. 2024. Text Embeddings by Weakly-Supervised Contrastive Pre-training. arXiv:2212.03533 [cs] doi:10.48550/arXiv.2212.03533 Comment: 17 pages, v2 fixes the SummEval numbers.
- [41] Jason Wei, Nguyen Karina, Hyung Won Chung, Yunxin Joy Jiao, Spencer Papay, Amelia Glaese, John Schulman, and William Fedus. 2024. Measuring Short-Form Factuality in Large Language Models. arXiv:2411.04368 [cs] doi:10.48550/arXiv.2411.04368 Comment: Blog post: <https://openai.com/index/introducing-simpleqa/>.
- [42] Wei Yang, Yuqing Xie, Aileen Lin, Xingyu Li, Luchen Tan, Kun Xiong, Ming Li, and Jimmy Lin. 2019. End-to-End Open-Domain Question Answering with BERTserini. In *Proceedings of the 2019 Conference of the North. 72–77*. arXiv:1902.01718 [cs] doi:10.18653/v1/N19-4013 Comment: Published in the Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics (Demonstrations).
- [43] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. 2018. HotpotQA: A Dataset for Diverse, Explainable Multi-hop Question Answering. arXiv:1809.09600 [cs] doi:10.48550/arXiv.1809.09600 Comment: EMNLP 2018 long paper. The first three authors contribute equally. Data, code, and blog posts available at <https://hotpotqa.github.io/>.
- [44] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. arXiv:2210.03629 [cs] doi:10.48550/arXiv.2210.03629 Comment: v3 is the ICLR camera ready version with some typos fixed. Project site with code: <https://react-lm.github.io>.
- [45] Zhehao Zhang, Jiaao Chen, and Diyi Yang. 2024. DARG: Dynamic Evaluation of Large Language Models via Adaptive Reasoning Graph. arXiv:2406.17271 [cs] doi:10.48550/arXiv.2406.17271
- [46] Kun Zhou, Yutao Zhu, Zhipeng Chen, Wentong Chen, Wayne Xin Zhao, Xu Chen, Yankai Lin, Ji-Rong Wen, and Jiawei Han. 2023. Don’t Make Your LLM an Evaluation Benchmark Cheater. arXiv:2311.01964 [cs] doi:10.48550/arXiv.2311.01964 Comment: 11 pages.
- [47] Kaijie Zhu, Jiaao Chen, Jindong Wang, Neil Zhenqiang Gong, Diyi Yang, and Xing Xie. 2024. DyVal: Dynamic Evaluation of Large Language Models for Reasoning Tasks. arXiv:2309.17167 [cs] doi:10.48550/arXiv.2309.17167 Comment: ICLR 2024 spotlight; 38 pages; code is at aka.ms/dyval.
- [48] Kaijie Zhu, Jindong Wang, Qinlin Zhao, Ruochen Xu, and Xing Xie. 2024. Dynamic Evaluation of Large Language Models by Meta Probing Agents. arXiv:2402.14865 [cs] doi:10.48550/arXiv.2402.14865 Comment: International Conference on Machine Learning (ICML) 2024.
- [49] Kaijie Zhu, Jindong Wang, Qinlin Zhao, Ruochen Xu, and Xing Xie. 2024. DyVal 2: Dynamic Evaluation of Large Language Models by Meta Probing Agents. arXiv:2402.14865 [cs] doi:10.48550/arXiv.2402.14865 Comment: Technical report; 20 pages.

Appendix

Claim Extraction Prompt

You are given a passage of text. Your task is to extract one-sentence standalone factual claims from the passage.

Each claim must:

- Be self-contained (i.e., understandable on its own without requiring coreference resolution)
- Not rely on ambiguous references like “this”, “that”, “they”, etc.
- Be phrased as a complete declarative sentence
- Be verifiable (e.g., “Riot Games has no plans for League of Legends 2.” is acceptable; “They have no plans for League of Legends 2” is not)

For each claim, also provide the exact verbatim supporting text span from the original text that led you to extract the claim.

Output JSON Keys:

- "claim1": standalone claim
- "supporting_text_span1": verbatim supporting text span from original
- "claim2": standalone claim
- "supporting_text_span2": verbatim supporting text span from original
- ...

Task:

Here is the input text: {text}

Temporal-Sequence Multi-Hop QA Prompt

You are generating {n_pairs} **temporal-sequence** multi-hop QA pairs.

Buckets of claims (grouped by doc_id) are below:

```
{json.dumps(buckets, indent=2)}
```

Rules for each QA pair:

- Combine claims from at least {min_buckets} **different buckets**.
- Each chosen claim must mention a year, date, or explicit timeframe.
- Ask for an ordering or an interval (e.g., “How many years after ...?”).
- Every referenced fact must be essential for the answer.
- The question should not contain the answer or steps to get the answer.
- The question should not refer to the documents, they should be general.
- List every claim you used by its doc_id and claim_id.
- Provide one concise answer (year, number of years, earlier event, etc.).

Output format (JSON list):

- "used_claims": list of {{doc_id, claim_id, claim}}
- "question": string
- "answer": string

Return only the JSON—no extra commentary.

Comparison-Style Multi-Hop QA Prompt

You are generating {n_pairs} **comparison-style** multi-hop QA pairs.

Buckets of claims (grouped by doc_id) are below:

```
{json.dumps(buckets, indent=2)}
```

Rules for each QA pair:

- Use claims from at least {min_buckets} **different buckets**.
- Frame the question so the solver must contrast values (higher, lower, difference, ratio, earlier, later).
- Keep wording tight—only include values essential for the comparison.
- The question should not contain the answer or steps to get the answer.
- The question should not refer to the documents, they should be general.
- List every claim you used by its doc_id and claim_id.
- Provide a single, concise answer (number, entity, etc.).

Output format (JSON list):

- "used_claims": list of {{doc_id, claim_id, claim}}
- "question": string
- "answer": string

Return only the JSON—no extra commentary.

Conjunction-Style Multi-Hop QA Prompt

You are generating {n_pairs} **conjunction-style** multi-hop QA pairs.

Buckets of claims (grouped by doc_id) are below:

```
{json.dumps(buckets, indent=2)}
```

Rules for each QA pair:

- Combine facts from at least {min_buckets} **different buckets** (can use more).
- The answer must change if any referenced claim were false (logical AND).
- No trivia—every fact must be essential.
- The question should not contain the answer or steps to get the answer.
- The question should not refer to the documents, they should be general.
- List every claim you used by its doc_id and claim_id.
- Provide a single, concise answer (entity, number, date, etc.).

Output format (JSON list):

- "used_claims": list of {{doc_id, claim_id, claim}}
- "question": string
- "answer": string

Return only the JSON—no extra commentary.

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009